

```
from google.colab import files

uploaded = files.upload()

Choose Files gsearch_jobs.csv
• gsearch_jobs.csv(text/csv) - 172727065 bytes, last modified: 11/26/2023 - 100% done
Saving gsearch_jobs.csv to gsearch_jobs.csv
```

```
import pandas as pd

# Assuming the file name is 'gsearch_jobs.csv'
file_name = 'gsearch_jobs.csv'

# Load data into Pandas DataFrame
df = pd.read_csv(file_name)

# Display the first few rows of the DataFrame
df.head()
```

Unnamed: 0	index	title	company_name	location	via	description	extensions
0	0	0	Data Analyst	Meta	Anywhere	via LinkedIn	In the intersection of compliance and analytic... [15 hours ago, '101K–143K a year', 'Work fro... eyJqb2JfdGI0bGUiO
1	1	1	Data Analyst	ATC	United States	via LinkedIn	Job Title: Entry Level Business Analyst / Prod... [12 hours ago, 'Full-time', 'Health insurance'] eyJqb2JfdGI0bGUiO
2	2	2	Aeronautical Data Analyst	Garmin International, Inc.	Olathe, KS	via Indeed	Overview:\n\nWe are seeking a full-time...\nAe... [18 hours ago, 'Full-time'] eyJqb2JfdGI0bGUiOiJE
3	3	3	Data Analyst - Consumer Goods - Contract to Hire	Upwork	Anywhere	via Upwork	Enthusiastic Data Analyst for processing sales... [12 hours ago, '15–25 an hour', 'Work from h... eyJqb2JfdGI0bGUiOiJf
4	4	4	Data Analyst Workforce Management	Krispy Kreme	United States	via LinkedIn	Overview of Position\n\nThis position will be ... [7 hours ago, '90K–110K a year', 'Contractor'] eyJqb2JfdGI0bGUiOi.

5 rows × 27 columns

```
df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 36051 entries, 0 to 36050
Data columns (total 27 columns):
#   Column              Non-Null Count  Dtype
---  -
0   Unnamed: 0          36051 non-null  int64
1   index               36051 non-null  int64
2   title               36051 non-null  object
3   company_name        36051 non-null  object
4   location             36025 non-null  object
5   via                  36042 non-null  object
```

```
6  description      36051 non-null object
7  extensions       36051 non-null object
8  job_id           36051 non-null object
9  thumbnail        20443 non-null object
10 posted_at        36051 non-null object
11 schedule_type    35874 non-null object
12 work_from_home   16025 non-null object
13 salary           6309 non-null object
14 search_term       36051 non-null object
15 date_time        36051 non-null object
16 search_location   36051 non-null object
17 commute_time     0 non-null float64
18 salary_pay        6309 non-null object
19 salary_rate       6309 non-null object
20 salary_avg        6309 non-null float64
21 salary_min        5952 non-null float64
22 salary_max        5952 non-null float64
23 salary_hourly     4167 non-null float64
24 salary_yearly     2129 non-null float64
25 salary_standardized 6309 non-null float64
26 description_tokens 36051 non-null object
dtypes: float64(7), int64(2), object(18)
memory usage: 7.4+ MB
```

df.describe()

	Unnamed: 0	index	commute_time	salary_avg	salary_min	salary_max	salary_hourly	salary_yearly
count	36051.000000	36051.000000	0.0	6309.000000	5952.000000	5952.000000	4167.000000	2129.000000
mean	18025.000000	1117.465729	NaN	34479.365127	29067.529400	40728.531630	42.456192	10205.000000
std	10407.171614	696.771915	NaN	51428.069516	43427.331132	60779.792689	22.815874	3076.000000
min	0.000000	0.000000	NaN	7.250000	8.000000	10.000000	7.250000	2920.000000
25%	9012.500000	530.000000	NaN	30.500000	18.330000	45.000000	25.000000	8375.000000
50%	18025.000000	1080.000000	NaN	57.500000	40.000000	75.000000	35.000000	9650.000000
75%	27037.500000	1643.000000	NaN	85000.000000	74752.500000	100000.000000	57.500000	11400.000000
max	36050.000000	3275.000000	NaN	288000.000000	230000.000000	346000.000000	300.000000	28800.000000

df.isnull().sum()

```
Unnamed: 0      0
index           0
title           0
company_name    0
location        26
via             9
description      0
extensions      0
job_id          0
thumbnail       15608
posted_at       0
schedule_type   177
work_from_home  20026
salary          29742
search_term     0
date_time       0
search_location 0
commute_time    36051
salary_pay      29742
salary_rate     29742
salary_avg      29742
salary_min      30099
```

```

salary_max      30099
salary_hourly    31884
salary_yearly    33922
salary_standardized 29742
description_tokens      0
dtype: int64

```

```
df['company_name'].unique()
```

```

array(['Meta', 'ATC', 'Garmin International, Inc.', ...,
      'Applied Memetics L.L.C', 'Global Enterprise Partners',
      'Techdash Telecom'], dtype=object)

```

```
df['title'].unique()
```

```

array(['Data Analyst', 'Aeronautical Data Analyst',
      'Data Analyst - Consumer Goods - Contract to Hire', ...,
      'Data Analyst - Business Intelligence',
      'COOP - Senior Data Analyst', 'Lead FP&A Analyst- Remote, US'],
      dtype=object)

```

```

import nltk
from nltk.corpus import stopwords
from nltk.tokenize import word_tokenize
nltk.download('stopwords')
nltk.download('punkt')

```

```

[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data] Unzipping corpora/stopwords.zip.
[nltk_data] Downloading package punkt to /root/nltk_data...
[nltk_data] Unzipping tokenizers/punkt.zip.
True

```

```
# Tokenize the job descriptions
```

```
df['description_tokens'] = df['description'].apply(lambda x: word_tokenize(str(x).lower()))
```

```
# Remove stop words and non-alphabetic tokens
```

```
stop_words = set(stopwords.words('english'))
```



```
df['description_tokens'] = df['description_tokens'].apply(lambda x: [word for word in x if word.isalpha() and word r
```

```
# List of known skills
```

```
known_skills = ["python", "java", "data analysis", "sql", "excel", "machine learning", "statistics", "r", "snowflake
```

```
# Display the first few rows to check the result
```

```
df[['description', 'description_tokens']].head()
```

	description	description_tokens	
0	In the intersection of compliance and analytic...	[intersection, compliance, analytics, seeking,...	
1	Job Title: Entry Level Business Analyst / Prod...	[job, title, entry, level, business, analyst, ...	
2	Overview:\n\nWe are seeking a full-time...\nAe...	[overview, seeking, aeronautical, data, analys...	
3	Enthusiastic Data Analyst for processing sales...	[enthusiastic, data, analyst, processing, sale...	
4	Overview of Position\n\nThis position will be ...	[overview, position, position, primary, person...	

```
# Create a new column in the DataFrame to store the extracted skills
```

```
df['extracted_skills'] = ""
```

```
# Loop through each row and tokenize the words to check for known skills
for index, row in df.iterrows():
    description_tokens = row['description_tokens']

    if description_tokens:
        extracted_skills = [skill for skill in known_skills if skill in description_tokens]
        df.at[index, 'extracted_skills'] = extracted_skills

from collections import Counter

# Combine all extracted skills into a single list
all_skills = [skill for skills_list in df['extracted_skills'] for skill in skills_list]

# Count the frequency of each skill
skill_counts = Counter(all_skills)


# Display the top 15 most common skills
top_skills = skill_counts.most_common(15)
print(top_skills)

[('sql', 18156), ('excel', 12061), ('python', 10350), ('statistics', 8590), ('r', 6550), ('snowflake', 2137), ('
```

```
import matplotlib.pyplot as plt
import seaborn as sns

# Convert skill counts to a DataFrame for easier plotting
skill_counts_df = pd.DataFrame(skill_counts.items(), columns=['Skill', 'Count'])

# Plot the top 15 skills
plt.figure(figsize=(12, 6))
sns.barplot(x='Count', y='Skill', data=skill_counts_df.head(15))
plt.title('Top Skills Mentioned in Job Descriptions')
plt.show()
```



```
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score
from sklearn.feature_extraction.text import CountVectorizer
```

```

# Create a subset (10% of the data)
subset_df = df.sample(frac=0.1, random_state=42)

# Check the shape of the subset to ensure it's a reasonable size
print("Subset shape:", subset_df.shape)

Subset shape: (3605, 28)

X = subset_df['description']
y = subset_df['title']

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

#Use CountVectorizer to convert 'skills' into numerical features
from sklearn.feature_extraction.text import CountVectorizer

vectorizer = CountVectorizer()
X_train_vec = vectorizer.fit_transform(X_train)
X_test_vec = vectorizer.transform(X_test)

# Train a RandomForestClassifier
clf = RandomForestClassifier(random_state=42)
clf.fit(X_train_vec, y_train)

RandomForestClassifier
RandomForestClassifier(random_state=42)

# Make predictions on the test set
y_pred = clf.predict(X_test_vec)

# Display actual labels and predicted labels side by side
predictions_comparison = pd.DataFrame({'Actual': y_test, 'Predicted': y_pred})
# Display the first 10 rows of the predictions_comparison DataFrame
N = 10
print(predictions_comparison.head(N))

```

	Actual \
35860	LEAD DATA ANALYST
29610	Data Scientist
14041	Senior Data Analyst (Transportation Strategy)
15731	Data Engineer, Planning & Analysis
17319	Systems Test Engineer - Test Data Analyst
34649	(USA) Analyst II, Merchandising Business Analy...
27998	Data Analyst with Claims
21553	Data Analyst
33150	Healthcare Analyst III
30552	Senior Financial Data Analyst (Remote) - Medic...
	Predicted
35860	LEAD DATA ANALYST
29610	Senior Data Analyst

```
14041 Senior Data Analyst
15731 Data Engineer, Planning & Analysis
17319 Data Analyst
34649 (USA) Analyst II, Merchandising Business Analy...
27998 Data Analyst
21553 Data Analyst at COVID TESTING LLC
33150 Analyst III - REMOTE
30552 Senior Financial Data Analyst (Remote) - Medic...
```

```
# Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)

Accuracy: 0.30235783633841884
```